

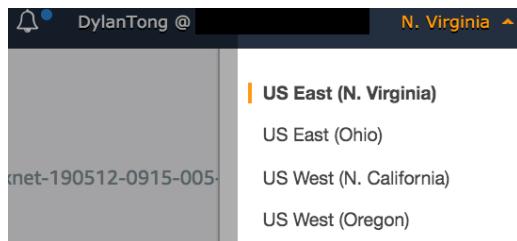
Amazon SageMaker Lab: Churn Predictive Analytics

Most Computer Vision projects involve *supervised* machine learning, which requires a training data set that contains ground truth information. For instance, an object detection model is typically trained on images that are annotated with bounding boxes around objects that you wish to detect, and corresponding labels that specify the class of an object—is it a person, a car, or a bird?

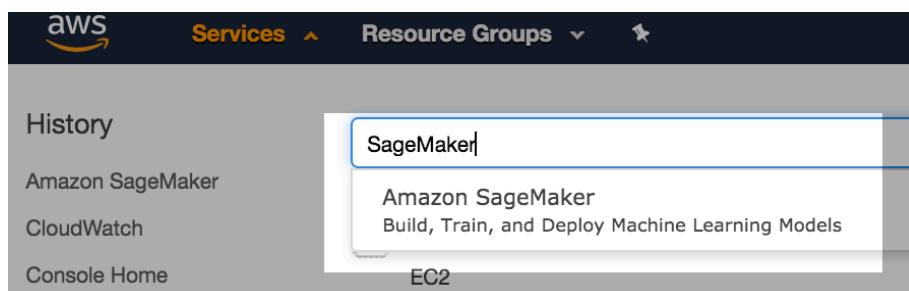
In this lab, we'll learn how to we can leverage Amazon SageMaker Ground Truth to create a training data set at scale. Challenges of annotating data at scale include dealing with the sheer volume of training data that needs to be produced on an ongoing basis and in a timely manner, managing the quality of the annotations, and integrating workflows with your machine learning tools and processes.

I. Prepare your Development Environment

1. Log into your AWS account and ensure you're in the right region designated for your workshop. The screenshot below indicates that I'm currently in us-east-1 (N. Virginia).



2. Navigate to the Amazon SageMaker console via the search bar.



3. Next, launch a managed Amazon SageMaker notebook instance. We're going to use this notebook instance to run a number of labs. In this lab, the notebook will be used to stage some raw data that we will annotate.

Switch to the **Notebook Instances** page by using the navigation menu on the left hand side of the console.

Name	Instance	Creation time	Status	Actions
cv-workshop-2	ml.c5d.4xlarge	May 13, 2019 21:55 UTC	InService	Open Jupyter Open JupyterLab
workshop-larger-notebook	ml.t2.large	May 13, 2019 21:00 UTC	Stopped	Start
dtong-cv-workshop	ml.t2.medium	May 12, 2019 22:25 UTC	InService	Open Jupyter Open JupyterLab
gluoncv-workshop	ml.p3.2xlarge	May 10, 2019 16:14	InService	Open Jupyter Open

4. Click on the **Create notebook instance** button.

5. Next, we configure our notebook instance by working through the launch wizard. First, provide a **name** for your notebook.

Utilize a unique prefix that you can remember, so you more easily find the resources that belong to you.

Notebook instance settings

Notebook instance name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

6. Select the ***ml.m5.4xlarge*** instance type.

Notebook instance type

ml.m5.4xlarge ▾

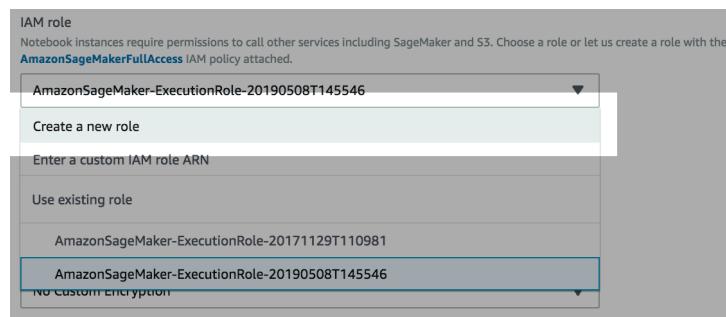
Elastic Inference [Learn more](#)

none ▾

► Additional configuration

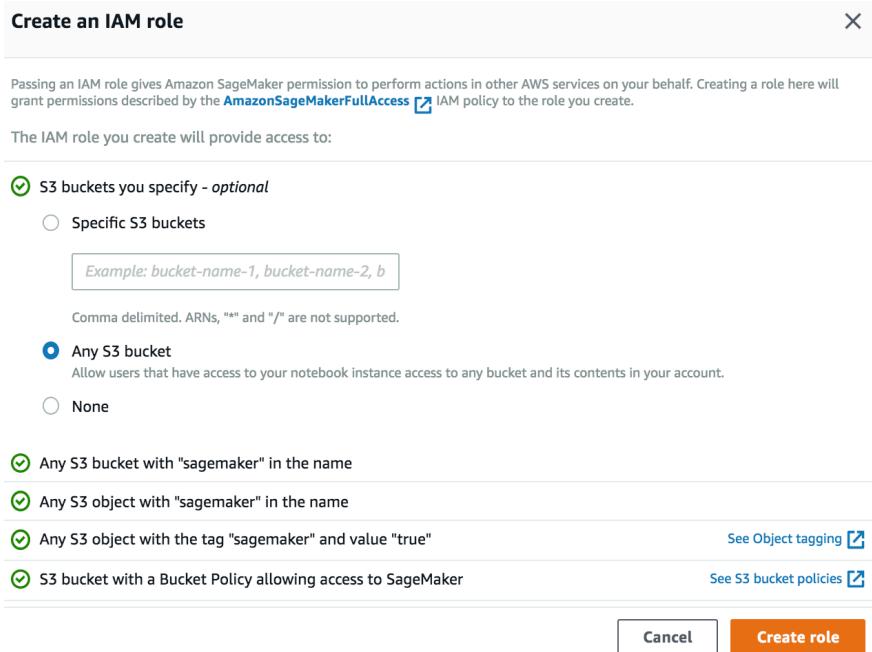
7. Your instance requires permissions to access data on S3, and execute SageMaker functionality required by this lab.

Create a new role for your notebook. You can select an existing role if you have already done this previously.



8. For the purpose of simplicity, select “**Any S3 bucket**.” If this wasn’t a lab, you should practice the principles of least privilege.

Click on the **Create role** button. Ensure that the role that you just created is selected in your notebook configurations.



9. Next, configure the Git integration. We're going to launch the notebook and clone the lab repository over to your notebook instance.

Select **“Clone a public Git repository to this notebook instance only.”**

Paste the following link into the text box under **“Git repository URL”**:
<https://github.com/dylan-tong-aws/aws-cv-jumpstarter>

The screenshot shows the 'Git repositories - optional' configuration section. It includes a 'Default repository' dropdown set to 'Clone a public Git repository to this notebook instance only' and a 'Git repository URL' text input field containing the URL 'https://github.com/dylan-tong-aws/aws-cv-jumpstarter'.

10. These basic configurations will suffice for the lab. In a production setting, you will likely want to launch this [notebook into a VPC](#) for better network security. [Life-cycle configurations](#) also come in handy if you like to automatically bootstrap your notebook instances with packages that aren't already pre-installed by default.

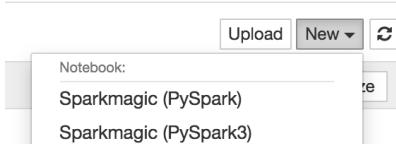
Click on “Create notebook instance” to launch your instance.



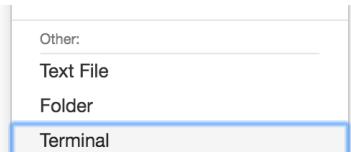
11. It will take about 5 minutes before your notebook is **InService**. Once it is, click on the “Open Jupyter” link.



12. Select the “New” drop down on the right hand side of the Jupyter admin console.



Scroll to the bottom and select **Terminal**.



13. We’re going to create an S3 bucket from the terminal. The AWS CLI has been pre-installed, and it inherits the IAM permissions of the role that you attached to the instance.

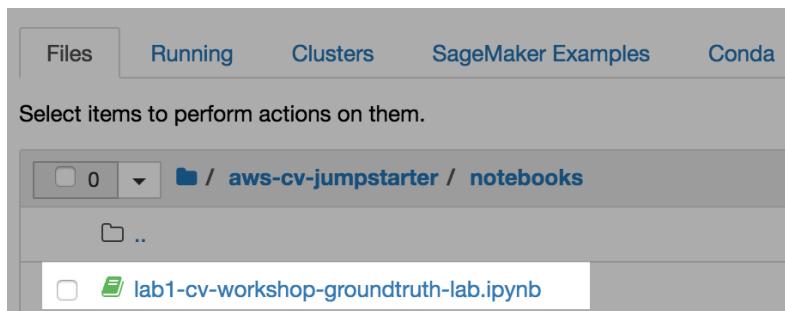
Run the following command and replace the parts that are **high-lighted in red** with appropriate values. First, your bucket needs a **unique name**. Secondly, you need to create your bucket in the **same region as your notebook instance**. The example below will create the bucket in Oregon (us-west-2).

```
aws s3api create-bucket --bucket dtong-cv-jumpstarter-workshop --region us-west-2 --create-bucket-configuration LocationConstraint=us-west-2
```

The output should look like the following:

```
sh-4.2$ aws s3api create-bucket --bucket dtong-cv-jumpstarter-workshop --region us-west-2 --create-bu  
cket-configuration LocationConstraint=us-west-2  
{  
    "Location": "http://dtong-cv-jumpstarter-workshop.s3.amazonaws.com/"  
}  
sh-4.2$
```

14. Return to the Jupyter admin console, and launch the Jupyter notebook for lab1 by clicking on the notebook as shown below:



15. Update the BUCKET variable with the name of the S3 bucket that you created previously from the terminal.

Replace << YOUR S3 BUCKET NAME >> with the name of your bucket.

```
BUCKET = '<< YOUR S3 BUCKET NAME >>'  
S3_PREFIX = 'ground-truth-lab' # Any valid S3 prefix.
```

For instance, the following is the name of the bucket that I created.

```
BUCKET = 'dtong-cv-jumpstarter-workshop'  
S3_PREFIX = 'ground-truth-lab' # Any valid S3 prefix.
```

16. Run through all the cells that you created. This notebook will transfer raw data over to your S3 bucket.

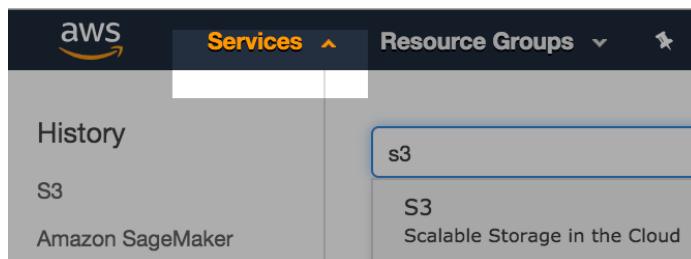


You should see a final output like the following. The process copies over 10 images, which we will annotate in the following steps.

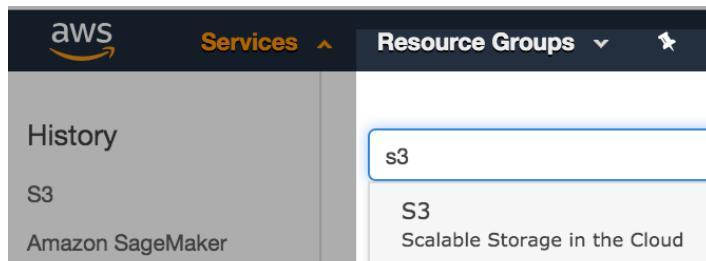
```
bbox_labels_600_hie 100%[=====] 84.27K --.KB/s in 0.07s
2019-05-12 22:35:32 (1.25 MB/s) - 'bbox_labels_600_hierarchy.json.l' saved [86291/86291]

Copying image 0 / 10
Done!
```

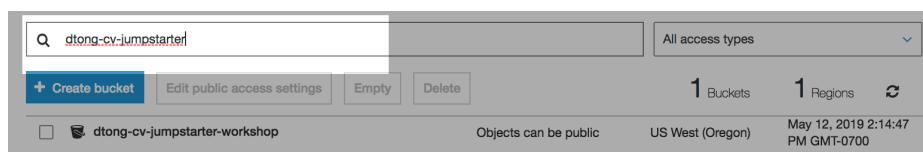
-
17. Let's validate what we've copied over into our S3 bucket. **Switch over to the S3 console.**



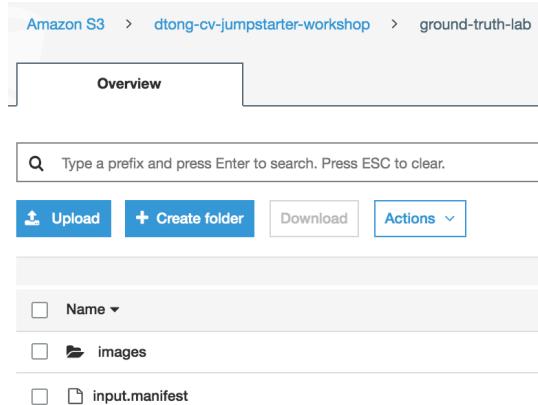
Use the search bar to link over to the S3 console.



18. **Select your bucket.** Use the search bar to filter the bucket list.



19. The contents of the bucket should appear as follows:



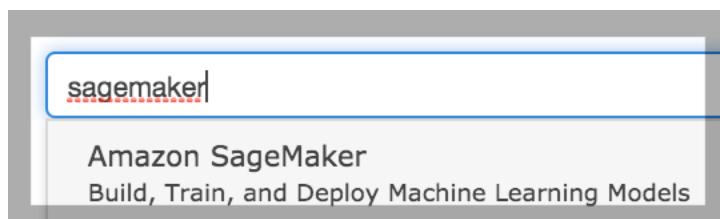
The images folder contains 10 images. If you download the input manifest, the contents should be like the following:

```
...
{"source-ref": "s3://dtong-cv-jumpstarter-workshop/ground-truth-
lab/images/000062a39995e348.jpg"}
{"source-ref": "s3://dtong-cv-jumpstarter-workshop/ground-truth-
lab/images/000411001ff7dd4f.jpg"}
...
...
```

The manifest file will be used by Amazon SageMaker GroundTruth to find the images that you wish to annotate. **Take note of the full path to this file.**

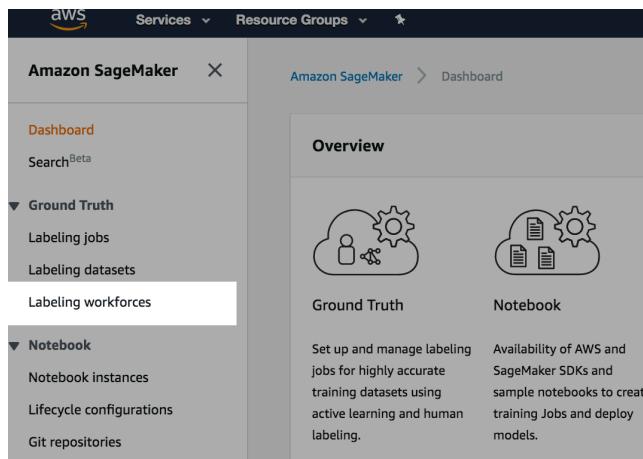
II. Prepare Your Labeling Workforce

20. Return to the **SageMaker** console.



21. Next, we're going to configure our a **Labeling workforce** to perform annotation on our behalf.

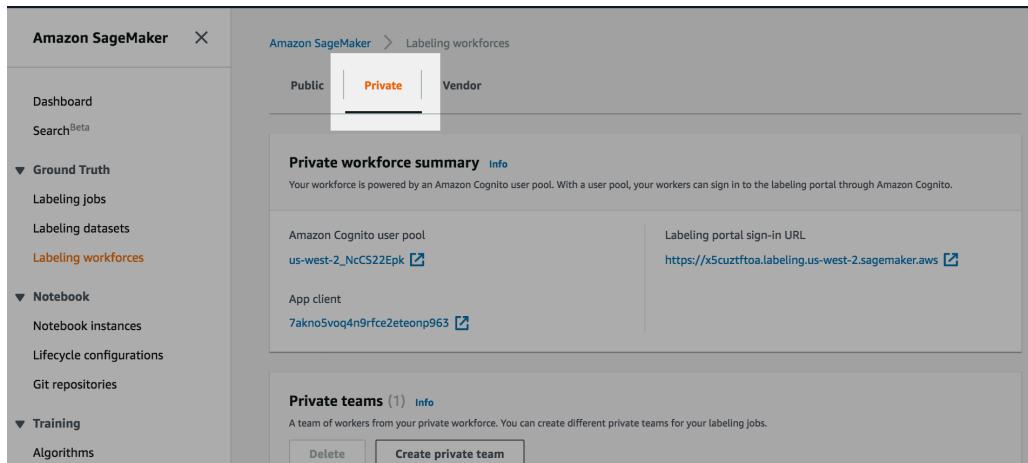
Select “**Labeling workforce**” from the navigation pane on the left hand side of the console.



22. SageMaker GroundTruth allows you to register labeling workforces from Mechanical Turk ([Public](#)), 3rd-party vendors registered in the AWS Marketplace ([Vendor](#)) as well as from our organization ([Private](#)).

We're going to create a private workforce in this lab consisting of yourself for the purpose of demonstration. This should be sufficient to give you a sense of how the other options would work in practice.

Select the **Private** sub tab from the **Labeling workforce** page.



23. Select **Create private team**.

The screenshot shows the Amazon SageMaker console with the 'Private' tab selected. On the left, there's a sidebar with options like Dashboard, Search (Beta), Ground Truth, Labeling jobs, Labeling datasets, Labeling workforces (which is highlighted in orange), Notebook, Training, and Algorithms. The main area shows a 'Private workforce summary' section with details about an Amazon Cognito user pool and an app client. At the bottom right of this section, there's a 'Create private team' button.

24. Enter a unique name for your team.

The screenshot shows the 'Team details' section. It has a 'Team name' field containing 'TeamDylan'. Below the field, there's a note: 'Maximum of 63 alphanumeric characters. Can include hyphens, but not spaces. Must be unique within your account in an AWS Region.'

25. Select “Invite new workers by email.”

Fill out the form as shown below with your email address in place of dylatong@amazon.com. You need to provide an email address that you have access to.

The screenshot shows the 'Email addresses' section. It has a text input field containing 'dylatong@amazon.com'. Below the input field, there's a note: 'Use a comma between addresses. You can add up to 50 workers.'

Below the email input, there are sections for 'Organization name' (containing 'AWS') and 'Contact email' (containing 'dylatong@amazon.com'). There are also 'Invite new workers by email' and 'Import workers from existing user groups' radio buttons, with 'Invite new workers by email' selected.

26. Leave the rest as default. You’re done with creating your private team. Click on the “Create private team” button.



III. Define your Labeling Job

27. Select **Labeling jobs** from the navigation panel on the left.

Name	Status	Task type	Labeled objects/total	Creation time
dtong-birds-labeling-job	Complete	Bounding box	10 / 10	May 12, 2019, 11:24 PM UTC

Click on the “Create labeling job” button.

28. Complete the labeling job wizard. First, provide a **unique name** for your labeling job—one that you can remember.

Job overview

Job name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. [View account in an AWS Region](#).

Provide the **full S3 URI** to the **input.manifest** file in your S3 bucket.

Input dataset location [Info](#)
 Provide a path to the S3 location where your manifest file is stored. To find a path, go to [Amazon S3](#)

[Create manifest file](#), if you don't have one

s3://dtong-cv-jumpstarter-workshop/ground-truth-lab/input.manifest

The bucket and dataset objects must be in the us-west-2 region

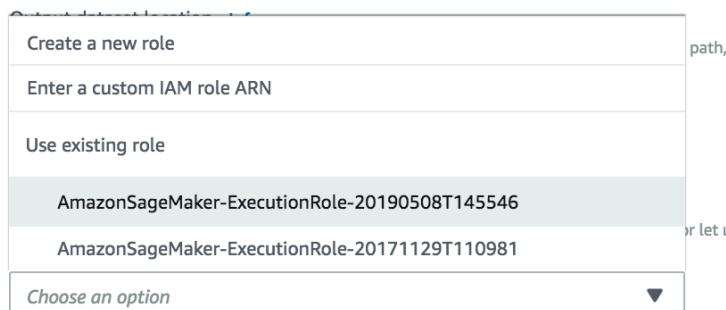
Specify an output path in your S3 bucket like the one below. Take note of this path as this is where the annotations created through SageMaker Ground Truth will be outputted.

Output dataset location [Info](#)
 Provide a path to the S3 location where you want your labeled dataset to be stored. To find a path, go to [Amazon S3](#)

s3://dtong-cv-jumpstarter-workshop/birds/train|

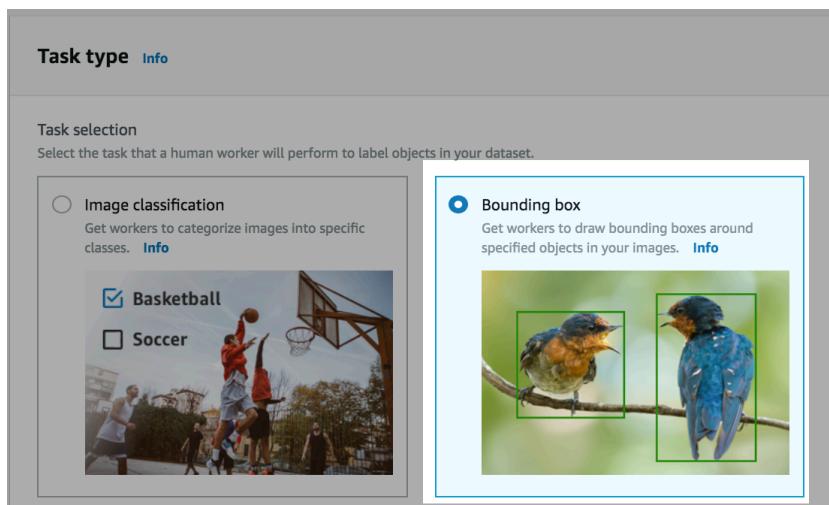
The bucket and dataset objects must be in the us-west-2 region

Select the SageMaker execution role that you created when you launched your notebook. The role will appear under “**Use existing role**.”

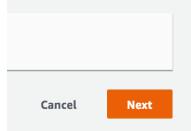


This role provides sufficient permissions for SageMaker Ground Truth. You should practice the least privileges principles outside of this lab. The role that you've created is overly permissive.

29. Under Task type, select **Bounding box**.



Click on Next.



30. Next, we'll assign this job to a workforce. There are a few options. Selecting **public** would assign the job to a Mechanical Turk workforce.

Workers [Info](#)

Worker types

Public
An on-demand 24/7 workforce of over 500,000 independent contractors worldwide powered by Amazon Mechanical Turk.

Private
A team of workers that you have sourced yourself, including your own employees or contractors for handling data that needs to stay within your organization.

Vendor managed
A curated list of third party vendors that specialize in providing data labeling services, available via the AWS Marketplace.

Price per task
We recommend you choose a price consistent with the approximate time it takes to complete a task. We have provided time estimates for each price as guideline to help you decide how you want to price your task.

\$0.036 [▼](#)
Time estimate: 8 secs - 10 s...

- The dataset does not contain adult content. [Info](#)
- I understand that my dataset will be viewed by the Amazon Mechanical Turk public workforce and I acknowledge that my dataset does not contain personally identifiable information (PII). [Info](#)

Selecting **Vendor managed** will allow you to assign the job to a 3rd-party vendor from the AWS Marketplace, which you have subscribed to. Billing needs to be setup in order for you to subscribe to a vendor.

Workers [Info](#)

Worker types

Public
An on-demand 24/7 workforce of over 500,000 independent contractors worldwide powered by Amazon Mechanical Turk.

Private
A team of workers that you have sourced yourself, including your own employees or contractors for handling data that needs to stay within your organization.

Vendor managed
A curated list of third party vendors that specialize in providing data labeling services, available via the AWS Marketplace.

Vendor managed data labeling services
The vendor managed labeling workforces enable you to delegate your labeling job to third parties who will deliver the labels.

Subscribed data labeling services
Before you choose a service, find and subscribe to vendor managed data labeling services on [AWS Marketplace](#)

[Choose a service](#) [▼](#)

I understand that when I select a vendor managed workforce for this data labeling job, my data will be shown to workers employed by the vendor as part of the labeling tasks.

You can navigate to the **AWS Marketplace** from the **shortcut** displayed below:

Workers Info

Worker types

- Public**
An on-demand 24/7 workforce of over 500,000 independent contractors worldwide powered by Amazon Mechanical Turk.
- Private**
A team of workers that you have sourced yourself, including your own employees or contractors for handling data that needs to stay within your organization.
- Vendor managed**
A curated list of third party vendors that specialize in providing data labeling services, available via the AWS Marketplace.

Vendor managed data labeling services
The vendor managed labeling workforces enable you to delegate your labeling job to third parties who will deliver the labels.

Subscribed data labeling services
Before you choose a service, find and subscribe to vendor managed data labeling services on [AWS Marketplace](#).

Choose a service

I understand that when I select a vendor managed workforce for this data labeling job, my data will be shown to workers employed by the vendor as part of the labeling tasks.

You can browse and subscribe to vendors from the AWS Marketplace as shown below.

aws marketplace

Categories ▾ Delivery Methods ▾ Solutions ▾ Migration Mapping Assistant Your Saved List Partners Sell in AWS Marketplace Amazon Web Services Home Help

Data Labeling Services (10 results) showing 1 - 10

iMerit

Data Labeling Services by iMerit
Machine Learning | Version 1 | Sold by iMerit, Inc.
iMerit delivers enterprise-grade data labeling and enrichment services in ML and computer vision across myriad use cases and verticals, so you can get the best out of your algorithms. Our full-time inhouse, India based staff of over 1600 data specialists provide secure, scalable and flexible expert...

Data Labeling Services by SmartOne, Inc.
Machine Learning | Version 1 | Sold by SmartOne, Inc.
SmartOne, Inc. is a data services company with more than six years of experience in the production of high-quality, high-volume datasets for AI and ML projects. We have more than 400 trained personnel with experience in computer vision, content moderation, text transcription, bounding box, and satellite...

Categories
All Categories
Machine Learning
Data Labeling Services
Filters
Clear all filters
Vendors
Vivetic (4)
SmartOne, Inc. (2)
iMERIT INC. (2)
SINFOSY GmbH (1)
STARTEK, Inc (1)
Software Pricing Plans
By Units (10)
Delivery Method
Clear

31. Select **Private** under **Worker types**.

Workers Info

Worker types

- Public**
An on-demand 24/7 workforce of over 500,000 independent contractors worldwide powered by Amazon Mechanical Turk.
- Private**
A team of workers that you have sourced yourself, including your own employees or contractors for handling data that needs to stay within your organization.
- Vendor managed**
A curated list of third party vendors that specialize in providing data labeling services, available via the AWS Marketplace.

Private teams
Choose from the teams you created in the private workforce or if you need to create a new team, save your progress and go to Labeling workforces to create a new one.

Select a private team

► Additional configuration - optional
Automated data labeling, workers per dataset object

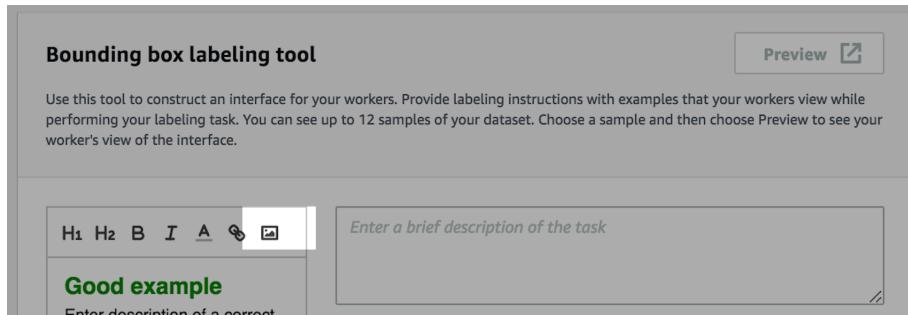
32. Select the **private workforce** that you created earlier from the **Private teams** drop down.

Private teams
Choose from the teams you created in the private workforce or if you need to create a new team, save your progress and go to Labeling workforces to create a new one.

Select a private team

 TeamDylan

33. We need to provide the labeling workforce with clear instructions. Providing a good and bad example is a good practice. **Select the icon** as shown below to link an image.



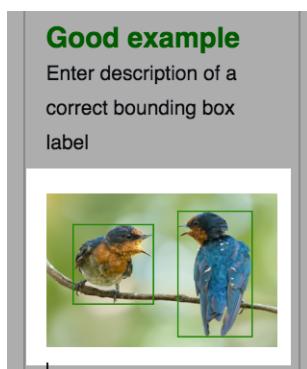
You can use the following image that is shared publicly from S3 through CloudFront:

<https://dvt7olt8euncl.cloudfront.net/41473cc4-ca5a-442f-9db9-cf116e59957f/src/images/bounding-box-good-example.png>

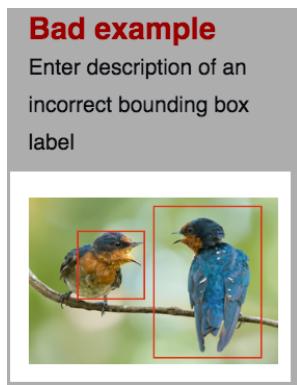
Copy the link into dialog box and select **Save**.



You should see the following image. Delete the placeholder, and leave this new image in its place.



Do the same with the following link to provide a “bad example:”
<https://dvt7olt8euncl.cloudfront.net/41473cc4-ca5a-442f-9db9-cf116e59957f/src/images/bounding-box-bad-example.png>



34. Provide instructions in the text box as shown below:

Bounding box labeling tool

Preview

Provide labeling instructions with examples below for workers. Workers will be viewing these instructions when they perform your task. You can add up to 10 labels for workers to choose from. See guidelines for creating high-quality instructions [🔗](#)

H₁ H₂ B I A
✖️ 🖼

Good example

Please create bounding boxes as illustrated by the images on the left. Thanks!

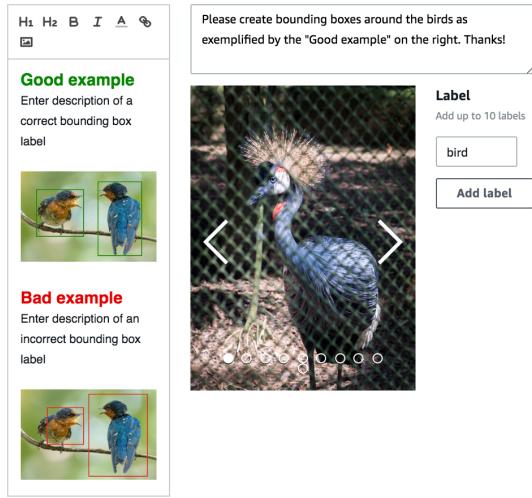
35. Enter “**bird**” into the **Label** textbox. In this lab we create a simple bird detector that doesn’t differentiate between different types of birds. In practice, you can add multiple labels for each object class you wish to classify.

Label

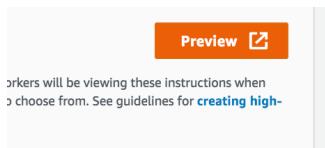
Add up to 10 labels

Add label

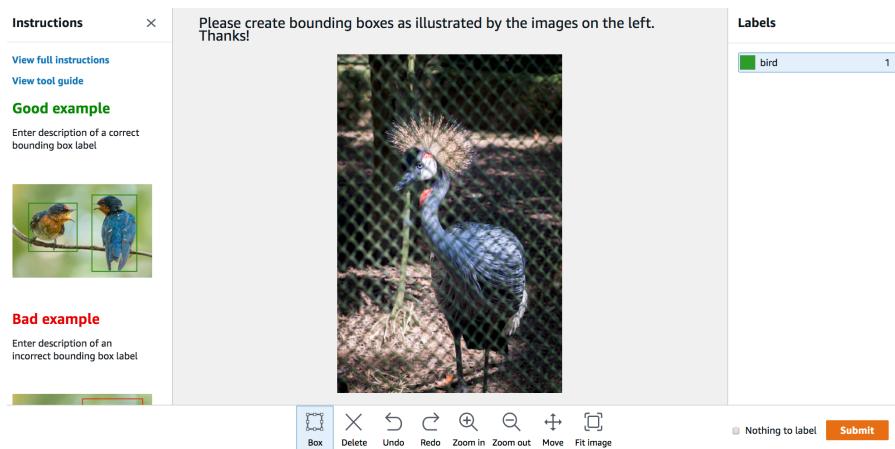
36. The completed template should appear like the following:



You can select the Preview button to get a glimpse of the interface that the labelers will see.

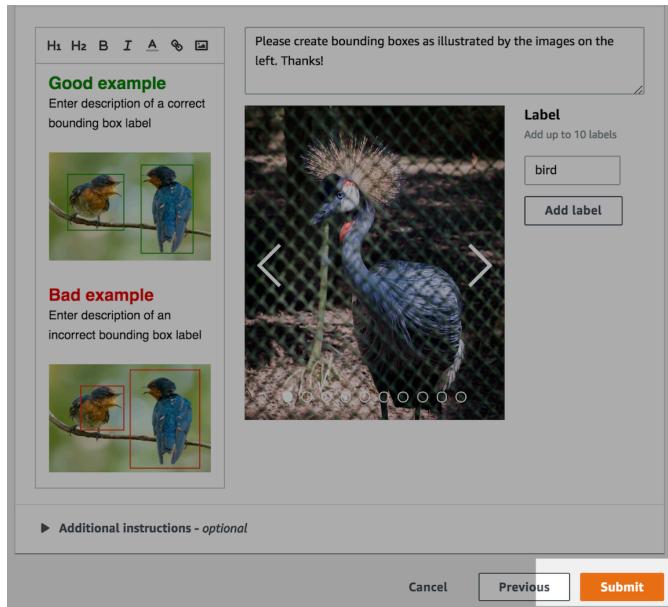


You should see something like this:



Close the window.

37. Click on the **Submit** button to complete the **Labeling job** description.



IV. Annotate your Data

38. Login to the email account, which you had an invitation sent out to earlier. The invitation should appear similar to the one below:

You're invited by AWS to work on a labeling project.

 no-reply@verificationemail.com <no-reply@verificationemail.com>
 Tong, Dylan
 Sunday, May 12, 2019 at 3:51 PM
[Show Details](#)

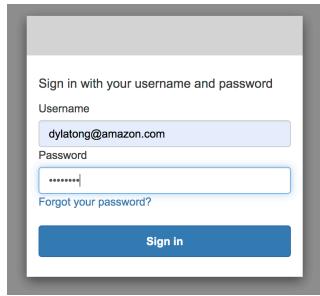
You're invited to work on a labeling project.

You will need this user name and temporary password to log in the first time.

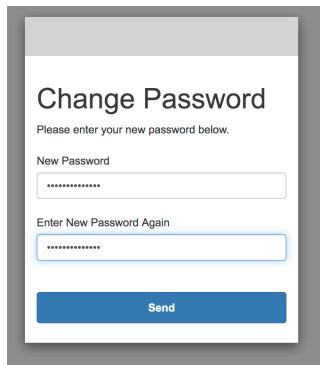
User name: dylatong@amazon.com
 Temporary password: **f@3RWFKH**
 Open the link below to log in:
<https://x5cuztftoa.labeling.us-west-2.sagemaker.aws>

After you log in with your temporary password, you are required to create a new one. If you have any questions, please contact dylatong@amazon.com.

Follow the instructions, and click on the link to login. Login with the temporary credentials that have been provided.



Change your password and remember it.

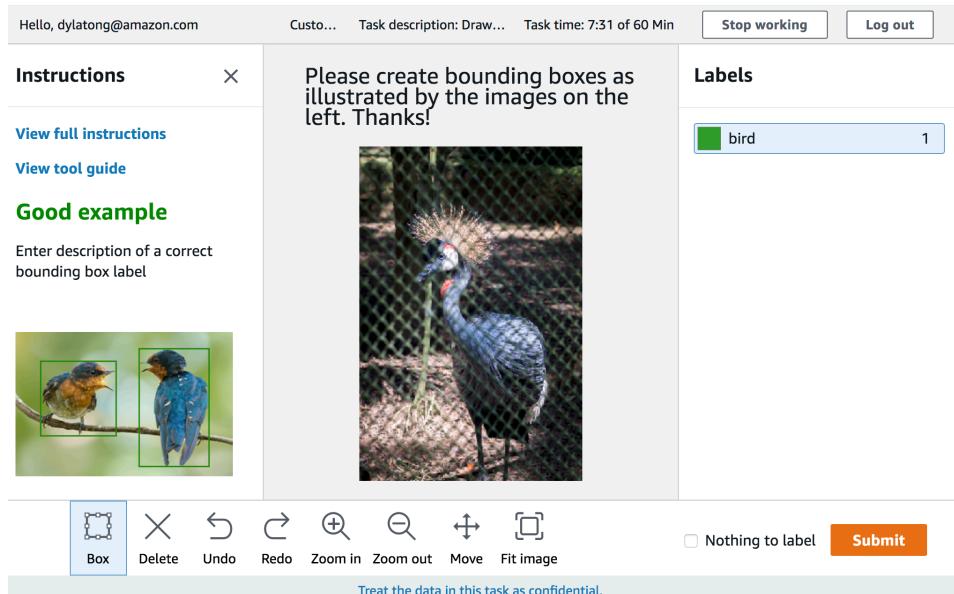


39. You should see the labeling job that you created earlier once you've logged into the portal. If not, try waiting a moment and refreshing your browser.

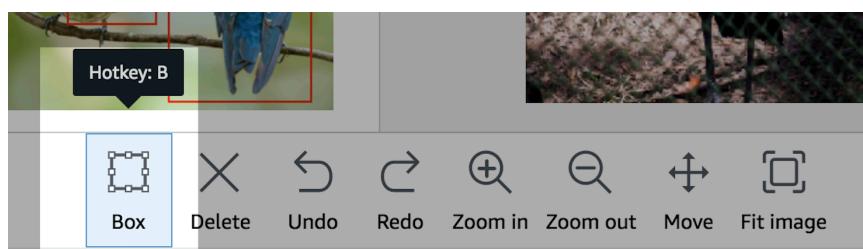
Select the job, and click **Start working**.

40. This labeling job consists of 10 images. You will need to annotate them, so that we can use them in the next lab.

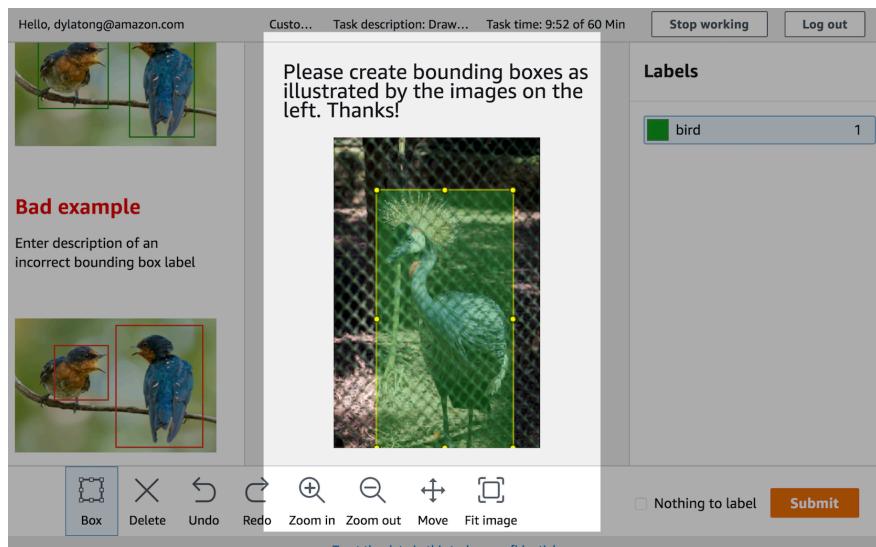
Create bounding boxes around each of the birds in the images.



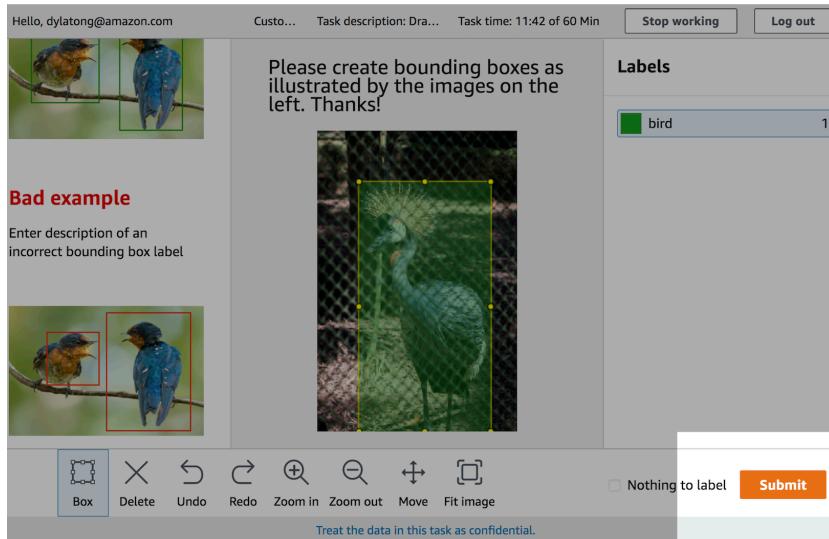
Select the **box** icon in the toolbar.



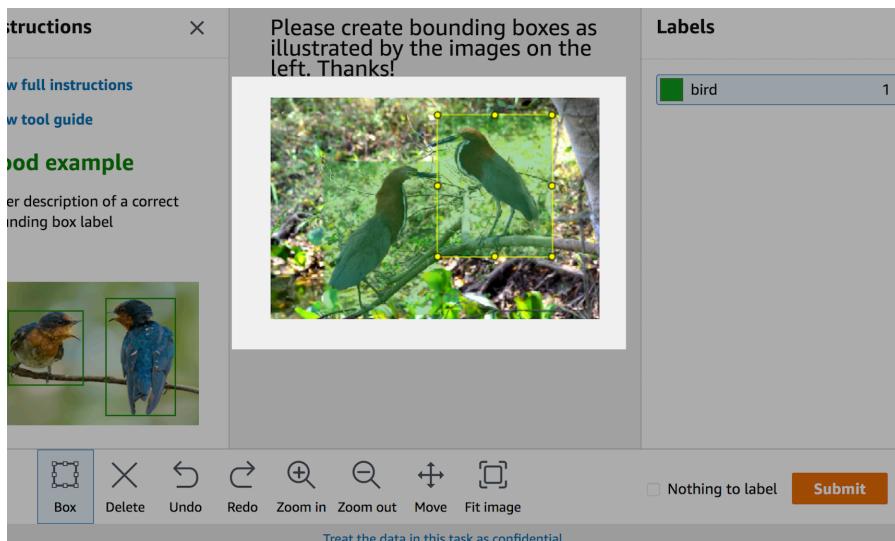
Click, drag and release to create the boxes.



Click on the **Submit** button once you're done with creating the bounding boxes.



Some of the images have more than one bird. Create boxes around all the birds like in the example below.



- Once you've annotated the 10 images, you'll be logged out of the annotation interface. You can view the status of your labeling job. Expect a short delay.

For manual (human) labeling jobs, SageMaker Ground Truth processes your annotations via [Annotation Consolidation](#) to improve the accuracy of labels. In the lab, you're the only worker in your private workforce, so Annotation Consolidation doesn't have any effect.

The screenshot shows the Amazon SageMaker Ground Truth interface. On the left, a sidebar lists 'Dashboard', 'Search Beta', 'Ground Truth' (which is expanded), 'Labeling jobs' (highlighted in orange), 'Labeling datasets', and 'Labeling workforces'. The main content area is titled 'Labeling jobs (1) Info'. It features a search bar and a table with columns: Name, Status, Task type, Labeled objects/total, and Creation time. The single row in the table is for 'dtong-birds-labeling-job', which is currently 'In progress' with a 'Bounding box' task type, 8/10 labeled objects, and created on May 12, 2019, at 11:24 PM UTC.

42. Once your labeling job shows up as **Complete**, click on the labeling job.

This screenshot shows the same 'Labeling jobs (1) Info' page as the previous one, but the job status has changed to 'Complete'. The table row for 'dtong-birds-labeling-job' now shows a green circular icon with a checkmark, indicating completion. The other details remain the same: 'Bounding box' task type, 10/10 labeled objects, and creation time of May 12, 2019, at 11:24 PM UTC.

Scroll down. Within the details of the job, you can preview your annotations.

This screenshot shows the 'Labeled dataset objects (10)' page. The sidebar on the left includes 'Dashboard', 'Search Beta', 'Ground Truth', 'Labeling jobs' (highlighted in orange), 'Labeling datasets', and 'Labeling workforces'. The main area displays four thumbnail images with blue bounding boxes around specific objects: a bird in a cage, two chickens, a lizard, and a bird on snow. Below each thumbnail is a file name: '000062a39995e348.jpg', '00411001ff7dd4f.jpg', '000062a39995e348.jpg', and '00411001ff7dd4f.jpg' respectively.

43. Click on the link “Output dataset location.”

This will take you to the location of the **output.manifest** file that was created by your labeling job. It contains meta-data for your image annotations.

The contents of the **output.manifest** should look something like this. Take note of the location of this file. In the next lab we'll run a training job, which can take this file as direct input without further modifications.

```
{
  "source-ref": "s3://dtong-cv-jumpstarter-workshop/ground-truth-lab/images/0027f99a032cca4a.jpg",
  "dtong-birds-labeling-job": {
    "annotations": [
      {
        "class_id": 0,
        "width": 578,
        "top": 147,
        "height": 382,
        "left": 139
      }
    ],
    "image_size": [
      {
        "width": 1024,
        "depth": 3,
        "height": 683
      }
    ]
  }
}
```